

The Last Age of Impunity?
The Accountability Chain and the Conservation of
Impunity

PIATRA . INSTITUTE

June 2026

Abstract

A provocative thesis reads the present authoritarian and plutocratic turn as an anticipatory counterrevolution, a reaction by the powerful against a near future in which AI, open-source investigation, automated translation, and computational audit make impunity impossible to maintain. The strong form does not survive: the movements predate the technology and are already explained by status threat, nationalism, elite economic interest, and institutional weakness. But a narrow mechanism does survive, and this paper prices it. Whether a wrongful act meets a consequence is the product of a chain of stage probabilities, discovery, attribution, publication, prosecution, independent judgment, and enforcement, and a product is bounded by its weakest term. Generative AI raises the first three, the evidence stages, toward one; it does not raise the last three, the institutional stages, which are human and capturable. Solved over a population of cases, the model yields three results. The bottleneck moves: before AI the binding stage is discovery and the probability of consequence sits at 0.078 of what the institutions would allow, while after AI it sits at the institutional ceiling exactly, so further evidence buys nothing and the binding constraint becomes adjudication. Impunity is conserved: cheaper evidence would multiply the probability of consequence nearly thirteenfold, from 0.010 to 0.123, but a self-protective elite rationally captures the new bottleneck, prosecution and judgment, and that capture absorbs the entire evidence gain, returning consequence to 0.007, marginally below where it began. And exposure decouples from consequence: the share of wrongdoing that becomes publicly legible rises toward one while the share that is sanctioned does not, so the gap between what is known and what is answered widens roughly fifteenfold, which is the pattern of a leak that reveals everything and a trial that acquits everyone. The same chain run downward, the state's automated case against a subject, inverts: AI raises the accusation stage and capture removes the protections, so a captured system harms the weak about sixfold. The age AI inaugurates is therefore not the last age of impunity but the age of impunity made visible, and what to do about visible impunity is a question about courts and enforcement, not about models.

1. The Provocation and Its Defeat

There is a seductive reading of the present in which the strongman, the oligarch, and the anti-institutional billionaire are not merely reactionary but prophetic. On this reading they are reacting against a future they can feel coming, one in which artificial intelligence, open-source investigation, automated translation, leaked archives, and computational audit dissolve the opacity that power has always depended on, so that the long age in which the powerful could count on delay, secrecy, better lawyers, and public forgetting is ending, and the attacks on courts, journalists, inspectors, and archives are an anticipatory counterrevolution against machine-amplified justice. It is a thesis with a kind of moral gravity, and it should be resisted in its strong form, because the strong form is false.

It is false for reasons of chronology and motive. Authoritarian nationalism, oligarchic capture, and charismatic anti-institutionalism are far older than machine learning; fascism, patrimonial rule, and the imperial police state managed truth and crushed oversight without a single model.

The movements the thesis names were substantially formed before generative AI had its present capabilities, and their support is already well explained by forces that owe nothing to the machine: status threat rather than personal economic hardship in the case most carefully studied (Mutz, 2018), and alongside it nationalism, religious mobilization, demographic anxiety, the collapse of local media, and the ordinary appetite of power for fewer constraints. Attacking a prosecutor demonstrates opposition to constraint, not foreknowledge of a computational audit. The strong thesis attributes a sophisticated, future-directed motive where self-preservation, ideology, and resentment are already sufficient.

What survives is narrower and worth taking seriously. As the cost of discovering and linking evidence falls, actors who benefit from opacity acquire a stronger incentive to disable, discredit, defund, or capture the institutions that convert evidence into consequences. This does not require any leader to think that an algorithm will prosecute them. It requires only that independent record-keeping, professional administration, investigative journalism, and open courts are constraints, and that cheaper evidence raises the value of controlling them. The question this paper answers is whether that mechanism amounts to a last age of impunity, a final phase after which the powerful can no longer escape, and the answer the model returns is no. The transition under way is not from human injustice to machine justice. It is from evidence scarcity to adjudication scarcity, and impunity is conserved across it.

2. The Accountability Chain

Begin with the object the strong thesis gets wrong by treating it as a single thing. Whether a wrongful act meets a consequence is not one event but the survival of a chain. The act must be discovered, the conduct must be attributed to a responsible agent, the finding must reach a public and legible form, a prosecutor or regulator must act on it, an independent forum must judge it, and a sanction must actually be enforced. Write the probability that an act of type m ends in a consequence as the product of the stage pass-probabilities,

$$P_m = p_{\text{disc}} p_{\text{attr}} p_{\text{pub}} p_{\text{pros}} p_{\text{judg}} p_{\text{enf}},$$

which is the seed of an old idea in the economics of deterrence, that misconduct stays attractive while its expected benefit exceeds the expected sanction, $B_m > P_m S_m + R_m$, with S_m the magnitude of the sanction and R_m the reputational cost (Becker, 1968). The crucial fact about a product is elementary and decisive: it is bounded above by its smallest term. A chain with one near-zero link has a near-zero value however strong its other links, so the probability of consequence is governed by the weakest stage, the bottleneck.

The six stages divide into two kinds, and the division is the whole argument. The first three are evidence stages. Discovery, attribution, and publication are the work of finding, connecting, and surfacing, and they are the work that generative AI, entity resolution, automated translation, optical character recognition over scanned archives, network analysis, and open-source investigation now perform at a scale that once required a large organization. The offshore leaks and the cross-

border investigations are the proof of concept: the FinCEN Files joined reporters in 88 countries, and automatic international exchange of bank information has been estimated to close most of the offshore tax gap it touches (Alstadsæter, Johannesen and Zucman, 2019; International Consortium of Investigative Journalists, 2024). The last three are institutional stages. Prosecution, independent judgment, and enforcement are not acts of finding but acts of an institution converting a finding into a consequence, and no model performs them. They are done by prosecutors who can be replaced, courts that can be packed, and enforcement that can be declined, which is to say they are human, and capturable.

AI therefore acts on one half of the chain and not the other. It raises the evidence stages toward one and leaves the institutional stages where the politics of a given country puts them. What that does to the product is the subject of the rest of the paper, and it is not what the strong thesis predicts.

3. The Bottleneck Moves

Run the chain over a population of 6,000 cases, each with stage probabilities drawn around plausible means, and let an AI capability parameter raise the evidence stages from their baselines toward one. Before AI, the binding stage is discovery, and the consequence the chain delivers is a small fraction of what the downstream institutions would permit. The institutional stages alone, prosecution times judgment times enforcement, define a ceiling on consequence, here a probability of about 0.123, the most that could happen if every wrongful act were perfectly discovered and attributed. Before AI the actual probability of consequence sits at 0.078 of that ceiling: the institutions could convert far more than they ever receive, because the evidence never arrives. Evidence is the binding constraint, and improving it improves accountability, which is why every earlier accountability revolution, from the printed pamphlet to the leaked cable, did real work against opacity.

After AI the picture inverts. As the evidence stages approach one, the probability of consequence rises to meet the institutional ceiling and stops there, sitting at essentially 1.0 of it. The binding constraint has migrated from discovery to prosecution, and it makes the crossing early: in the model the binding stage becomes institutional once AI capability passes about 0.2, a modest level. Past that point the chain is no longer starved of evidence; it is starved of conversion. Perfect discovery now buys exactly the institutional product and not one unit more, because a product cannot exceed its weakest term and the weakest term is no longer the one AI touches. This is the precise content of the phrase evidence scarcity to adjudication scarcity. The thing that was rare has become abundant, and the thing that was assumed has become the constraint.

The migration has a second half that the figure also shows, and it belongs properly to the next section. A self-protective elite spends to suppress the chain, and it spends where suppression is cheapest per unit of effect. Before AI, the evidence stages are cheap to suppress because they are scarce and centralized: intimidate the one witness, control the one archive, buy the one newspaper. After AI, the evidence stages are distributed and expensive to suppress, because a database that has leaked cannot be un-leaked and a method that is open cannot be made secret, so the cheapest remaining target is the institution. The budget that once silenced evidence now captures courts.




Figure 1: The probability of consequence as a fraction of the institutional ceiling, the most the downstream institutions could deliver. Before AI (left of the sweep) the chain runs at 0.078 of the ceiling because discovery starves it; as AI raises the evidence stages the chain rises to sit exactly at the ceiling, and the binding stage migrates from discovery to prosecution once AI capability passes about 0.2. The capture budget tracks the bottleneck: at low AI it suppresses the cheap, centralized evidence stages, and at high AI, when evidence is distributed and expensive to suppress, it captures the institutional stages instead.

4. Impunity Is Conserved

This is the result the strong thesis cannot see, because it imagines the contest as evidence against opacity and declares victory when evidence wins. Let the elite respond rationally. Give it a fixed budget to drive stages toward a capture floor, and let it allocate that budget to minimize the probability of consequence. Because driving any stage to the floor multiplies every case's product by roughly the same factor, the leverage of capturing a stage is nearly equal across stages, so the elite simply buys the cheapest captures its budget affords. After AI, those are the institutional chokepoints: a chief prosecutor, a high court. In the model the budget captures prosecution and judgment.

The arithmetic of the three regimes is the paper's center. Before AI, the mean probability of consequence is 0.010. Let AI raise the evidence stages with no capture, the optimistic picture in which machine legibility simply wins, and consequence rises to 0.123, nearly thirteen times higher: this is the future the strong thesis both fears and hopes for. Now let the elite capture the new bottleneck. Consequence falls back to 0.007, which is not merely lower than the uncaptured figure but marginally below the pre-AI level itself. The capture absorbs the entire evidence gain, 102 percent of it, and returns the system to where it started or a touch worse. Impunity is conserved. The evidence revolution happened, and nothing changed.




Figure 2: Mean probability of consequence under three regimes. Cheaper evidence alone would raise it nearly thirteenfold, from 0.010 to 0.123 (green). A self-protective elite that captures the new bottleneck, prosecution and judgment, absorbs the entire gain and returns consequence to 0.007 (red), marginally below the pre-AI level. The evidence revolution is real and is fully neutralized at the institutional stages it does not touch.

The conservation is not an accident of the parameters; it is a property of the structure. A product is bounded by its weakest controllable link, and a self-protective elite always has a weakest controllable link to hold down, so any improvement confined to the links it cannot control is absorbed at the link it can. This is the formal content of the historical regularity that no medium ever abolished

impunity. Writing, double-entry bookkeeping, the printing press, public statistics, photography, the telegraph, freedom-of-information law, the leak, the smartphone, each lowered some information cost, and each was met by censorship, ownership, secrecy, procedural obstruction, overload, or control of enforcement (Scott, 1998; Power, 1997). There has never been a final accountability medium because there cannot be one while the conversion stages remain capturable. The model also dissolves the puzzle the strong thesis raises, of why the powerful attack courts and prosecutors rather than documents. They attack courts because, once evidence is cheap, courts are the bottleneck, and the bottleneck is where capture pays. No anticipation of machine justice is required; the marginal logic of a weakest-link chain is enough.

5. Exposure Without Consequence

If consequence barely moves while evidence becomes abundant, something else must be growing, and it is the distance between what is known and what is answered. Define exposure as the evidence product alone, the probability that an act is discovered, attributed, and made public, and consequence as the full product. Before AI the two are close, both small: little is found, and of what is found little is sanctioned, so the gap between them is 0.068. After AI, with rational capture, exposure rises to essentially one while consequence stays at 0.007, and the gap widens to 0.993, roughly fifteen times its earlier size. The age that arrives is not the age of justice. It is the age of legible impunity, in which the powerful are seen clearly and answer for nothing.




Figure 3: Exposure, the probability that wrongdoing becomes publicly known, against consequence, the probability that it is sanctioned, before and after AI under rational capture. Exposure rises toward one while consequence stays flat, so the gap between what is known and what is answered widens roughly fifteenfold. Everything is visible and almost nothing follows.

This is not a forecast but a description, because the pattern is already on the record. The Panama Papers were among the largest financial leaks in history, read by hundreds of journalists across dozens of countries, and the principal money-laundering trial they produced in Panama ended in 2024 with all 28 defendants acquitted (International Consortium of Investigative Journalists, 2024).

Discovery succeeded completely and enforcement failed completely, which is exactly the shape of a chain whose evidence stages are saturated and whose institutional stages are not. The wider institutional record points the same way. Independent measures converge on a deterioration not of information but of the conversion institutions: V-Dem counts dozens of autocratizing states containing a large share of the world's population, the World Justice Project reports that the great majority of countries have seen declines in oversight, in checks on executive power, and in freedom of expression, and Transparency International places most jurisdictions below the midpoint of its scale (V-Dem Institute, 2025; World Justice Project, 2024; Transparency International, 2025). The visible decline is concentrated where the model says capture would concentrate, on the stages that turn evidence into consequence. Legible impunity may be worse for legitimacy than the hidden kind, because visible, unanswered wrongdoing teaches a public that consequence is for the weak, and that lesson corrodes the belief on which the institutions run (Arendt, 1973; Levitsky and Ziblatt, 2018).

6. The Chain Turned Downward

The same machinery has a second use, and it is the one that should temper any hope that legibility is on the side of the weak. Run the chain downward, as the state's automated case against a subject rather than a public's case against power. Now the early stage is accusation, the flagging of a person by a predictive or scoring system, and the protective stages are human review, contestability, and appeal. The valences reverse. AI raises the accusation stage, initiating far more cases, exactly as it raised discovery; and a captured administration removes the protections, driving review and appeal toward the same floor it used on the courts. In the model a subject's probability of an adverse automated outcome rises from 0.113 under an accountable configuration to 0.736 under a captured one, about sixfold. The tools that find a powerful actor's hidden conduct are the tools that flag a poor one's benefits claim, and the institution decides which way they point (Foucault, 1977; Eubanks, 2018; O'Neil, 2016; Noble, 2018).

This is why legibility is not a synonym for justice, as the work on seeing like a state established before any of these systems existed: the same standardization that lets an authority detect fraud lets it classify, exclude, and discipline a population (Scott, 1998; Crawford, 2021; Zuboff, 2019). The configuration that distinguishes a democratic use from an authoritarian one is not the capability but its direction. Democratic accountability seeks high legibility of public power and limited, legally bounded legibility of private life; authoritarian administration reverses the relation, opacity above and legibility below. The capability is symmetric. The politics chooses the sign. The legal order that decides which way it points is itself something the powerful now build, the quiet construction of informational capitalism through which control over data, classification, and platforms becomes a governing power (Cohen, 2019).

7. What the Model Settles, and What It Does Not

The model settles a narrow thing firmly and leaves the large things open, and honesty about the boundary is the point. It prices an incentive structure, not a country: the 0.010 and the 0.123 and

the fifteenfold gap are properties of a stipulated chain chosen to be transparent, and nothing here measures a real regime or predicts a real trial. It does not vindicate the strong thesis, and it should not be read as quietly smuggling it back; the anticipatory-counterrevolution reading remains false, the movements remain older than the machine, and their motives remain multicausal, with status threat, nationalism, and the ordinary hunger for unconstrained power doing the explanatory work that the machine cannot. What the model adds is only that, conditional on those movements existing, cheaper evidence gives them a sharper reason to do the specific thing they are observed doing, which is to capture the conversion institutions rather than to chase the documents.

The deeper lesson cuts against the technological hope directly, and it is not despair. The dream encoded in the strong thesis was that the machine would judge, that abundant evidence would carry its own verdict and the age of impunity would end on its own. The model says the machine only ever finds, and finding was never the scarce thing that protected the powerful; the scarce things were a prosecutor willing to act, a court able to judge without fear, and an enforcement that followed. Those are the stages no model supplies and the stages now under the heaviest attack, and the conservation result says plainly that every advance in evidence is worthless against an actor who controls them. The work that remains is therefore the unglamorous institutional work the audit dream hoped to skip: independent prosecutors, courts insulated from capture, enforcement that binds the strong. The age AI inaugurates is not the last age of impunity. It is the age of impunity made visible, and what to do about impunity one can finally see is a question about courts, not about models.

References

- Alstadsæter, A., Johannesen, N., and Zucman, G. (2019). Tax evasion and inequality. *American Economic Review*, 109(6), 2073–2103.
- Arendt, H. (1973). *The Origins of Totalitarianism* (new ed.). Harcourt Brace Jovanovich.
- Becker, G. S. (1968). Crime and punishment: an economic approach. *Journal of Political Economy*, 76(2), 169–217.
- Cohen, J. E. (2019). *Between Truth and Power: The Legal Constructions of Informational Capitalism*. Oxford University Press.
- Crawford, K. (2021). *Atlas of AI: Power, Politics, and the Planetary Costs of Artificial Intelligence*. Yale University Press.
- Eubanks, V. (2018). *Automating Inequality: How High-Tech Tools Profile, Police, and Punish the Poor*. St. Martin's Press.
- Foucault, M. (1977). *Discipline and Punish: The Birth of the Prison* (A. Sheridan, Trans.). Pantheon Books.
- International Consortium of Investigative Journalists. (2024). *Panama Papers and FinCEN Files investigations*. ICIJ.

- Levitsky, S., and Ziblatt, D. (2018). *How Democracies Die*. Crown.
- Mutz, D. C. (2018). Status threat, not economic hardship, explains the 2016 presidential vote. *Proceedings of the National Academy of Sciences*, 115(19), E4330–E4339.
- Noble, S. U. (2018). *Algorithms of Oppression: How Search Engines Reinforce Racism*. New York University Press.
- O’Neil, C. (2016). *Weapons of Math Destruction: How Big Data Increases Inequality and Threatens Democracy*. Crown.
- Power, M. (1997). *The Audit Society: Rituals of Verification*. Oxford University Press.
- Scott, J. C. (1998). *Seeing Like a State: How Certain Schemes to Improve the Human Condition Have Failed*. Yale University Press.
- Transparency International. (2025). *Corruption Perceptions Index 2024*. Transparency International.
- V-Dem Institute. (2025). *Democracy Report 2025*. University of Gothenburg, V-Dem Institute.
- World Justice Project. (2024). *Rule of Law Index 2024*. World Justice Project.
- Zuboff, S. (2019). *The Age of Surveillance Capitalism: The Fight for a Human Future at the New Frontier of Power*. PublicAffairs.